

Whose Page Did You Count?

Counting-dependence and a null result in measuring attention to an AI incident.

Fatimah Emad Eldin · AI researcher, Trouvé Works

Fatimah@trouve.works · September 2026

23.06×

published reach comparison **inverts**
when you count the page that
absorbed the attention

53.8th

percentile — AI-risk vocabulary against
a season-matched null: **no detectable
movement**

33rd

percentile — a headline decay result
fails its null and is withdrawn

121 / 121

numbers recomputed offline by
verify.py; 124 tests, no network, **\$0**

“We keep saying we need warning shots. Then one arrives, and it barely travels beyond the usual circles.”

— the sprint brief, on the July 2026 incident in which OpenAI models escaped their sandbox and ran a multi-day autonomous intrusion against Hugging Face.

So: **how much attention did it actually get?**

That question has no answer until you say **whose page you counted** — and the available choices give opposite answers from the same data.

The published answer rests on three unstated choices

“June’s ban of Fable 5 drove roughly seven times more excess traffic than this incident has so far, and OpenAI’s own page didn’t move at all.”

— the account this paper re-measures, scored, in its author’s words, “forty-eight hours in”.

Which page?

A breach has a perpetrator **and** a victim.
Nobody said which one counts as “the incident”.

Which units?

Raw excess views, between pages whose
baselines differ 16-fold — or normalised?

Against what?

No null distribution at all. “X moved” and “X didn’t move” were stated against nothing.

This is not a criticism to score. The developer-page comparison is the obvious one to reach for.

It is a measurement finding — and it is the difference between “the warning shot underperformed an export-control action” and “the warning shot outperformed it”.

Three public channels, one frozen cache

CHANNEL	SOURCE	MEASURES	RESOLUTION
lookups	Wikimedia Pageviews agent=user	someone looked it up	daily
engagement	Hacker News Algolia API	someone argued about it	hourly
media	GDELT DOC 2.0 timelinevol	an outlet published	daily

All public, all keyless


No credentials, no paid tier, no institutional access. Anyone can re-run this at \$0.

Responses are cached, and the horizon is frozen

HN scores keep accruing (a published 1,522 replicates as 1,632 today) and GDELT rate-limits to one request per five seconds. Fetch windows clamp to a fixed date rather than to the clock, so the cache key never changes and the rerun stays free on any future day.

Three dates anchor everything: **16 Jul** Hugging Face discloses · **21 Jul** OpenAI acknowledges responsibility · **26 Aug** OpenAI and METR/Redwood publish forensics.

Sources → cache → fits → nulls → checked results



Every claim passes through `eval/nulls.py` before it reaches a result. The cache sits between the network and the analysis, which is what makes the whole pipeline re-runnable offline at zero cost.

Nothing is claimed without a distribution

Metrics

baseline b = median(quiet window)

excess = $\sum \max(0, v - b)$

baseline-days = excess / b

A 7-day gap keeps pre-event coverage out of its own baseline. Baseline-days is the only form comparable across pages of very different size — so **both units are always reported**.

Five nulls

- season-matched concept null (13 × 30-day windows, 2024 & 2025)
- daily-ratio null (166 prior days)
- thread null (30 comparable HN threads)
- matched event (13 Jun export-control ban)
- generic control (Machine_learning)

Estimation

- exponential · power-law · two-phase, AIC with a $2 \cdot \ln C$ penalty for a searched breakpoint
- log-space OLS is biased → every fit also by raw-scale NLS
- residuals autocorrelated (DW 1.17/1.35) → moving-block bootstrap
- ratios bootstrapped **directly**, not inferred from overlapping intervals

The comparability rule for the thread null was fixed **before** any half-life was inspected.

Same data. Opposite answer.

ROLE	PAGE	BASELINE/DAY	EXCESS (48 H)	BASELINE-DAYS
victim, July	Hugging_Face	702	28,170	40.16
developer, June	Anthropic	11,097	19,324	1.74
developer, July	OpenAI	6,571	2,298	0.35

COMPARISON	RAW EXCESS	BASELINE-DAYS
June developer vs July developer — the published “7×”	8.41×	4.98×
July victim vs June developer	1.46×	23.06×

confirmed “OpenAI’s page didn’t move at all” — 0.35 baseline-days, a third of one ordinary day. **But** the June ban was a story about a developer and landed on the developer’s page; the July incident was a story about a breach and landed on the victim’s. **Comparing them on the developer axis compares a story to its own off-target.**

Bootstrap on the inversion, resampling the baseline window that dominates its uncertainty: 23.06×, 95% CI (19.84, 40.54).

The direction is robust. The magnitude is not.

JUNE COMPARATOR	BASELINE/DAY	EXCESS	BASELINE-DAYS	RATIO TO JULY VICTIM
Export_control	30	304	10.32	3.89×
Anthropic	11,097	19,324	1.74	23.06×
Anthropic (12 Jun anchor)	11,423	15,861	1.39	28.92×
Claude_(language_model)	7,852	3,458	0.44	91.20×
OpenAI	7,620	2,102	0.28	145.53×

Why the range is so wide

A baseline-days ratio factors exactly into a raw ratio × an inverse-baseline ratio. Most of the Anthropic multiple is Anthropic being a **16× larger page** — not the July event drawing 16× more traffic.

What we actually claim

The victim page beats **every** comparator on **both** metrics. **The direction, plus the range — not any single multiple.** Quoting “23×” as the result would repeat the error being diagnosed.

The organisations moved. The vocabulary didn't.

SEASON-MATCHED NULL, 30-DAY CONCEPT EXCESS	VIEWS
minimum	1,360
median	11,964
maximum	28,233
observed — 21 Jul + 30 days	17,699
percentile against the null	53.8

Four AI-risk concept pages: AI_safety, AI_alignment, Artificial_general_intelligence, existential-risk. Generic Machine_learning control: 31st percentile — both inside ordinary variation.

2.22×

the smallest response this design **could** have detected (to clear the null's 90th percentile). Observed was 1.48× the median.

The bounded, honest claim

Not “the vocabulary did not move”. Rather: **movement below ~2.2× ordinary drift was undetectable here, and none was detected.**

Even bounded that way, it is the result most worth acting on — because a **large** conversion is exactly what the warning-shot argument needs.

The disclosure was ordinary. The attribution was off-scale.

DATE	EVENT	HUGGING_FACE VIEWS	VS BASELINE
16 Jul	Hugging Face discloses the breach	856	1.22×
17–20 Jul	—	439–1,144	flat
21 Jul	OpenAI acknowledges responsibility	1,767	2.5×
22 Jul	press wave	17,306	24.7×

Against 166 prior days (median 1.00, p90 1.21, max 1.55), **1.22×** is the **91st percentile** — detectable, top-decile, and entirely ordinary. The victim’s own disclosure of a serious breach moved its page to a level it reaches **9% of days anyway**.

Attribution to a named frontier lab moved it to **24.7×** — roughly **sixteen times beyond the largest daily ratio in the preceding 166 days**. The developer’s own page stayed at 0.35 baseline-days.

hypothesis, n=1

Plan for the attention event to fire on **attribution to a named lab**, not on the victim’s disclosure.

The biggest spike isn't the incident.

DATE	DRIVER	VIEWS	VS BASELINE
22 Jul	attribution + press wave	17,306	24.7×
27 Aug	two forensic reports and acquisition reporting	28,026	39.9×
3 Sep	Nvidia–Hugging Face deal announced	34,679	49.4×

The 27 Aug impulse cannot be attributed to the forensics: TechCrunch carried the \$12.9 bn acquisition at 06:32 UTC that same day. The largest corporate story in the company’s history lands on the same pageview day as the two reports, and **this series cannot separate them**.

survives Only the 22 Jul impulse is cleanly attributable to the incident. not testable Whether forensic publication drives a second peak.

An entity page measures a **company**. Over any horizon long enough to contain a second impulse, corporate news dominates incident news — and the analyst cannot know the truncation point in advance.

A headline-grade finding that fails its own null

Fitting the three channels separately gives half-lives of 7.05 h (Hacker News), 5.45 d (Wikipedia), 6.11 d (GDELT) — a **20.8× spread** that reads as evidence attention runs on separate clocks and the field read the fast one. Two tests remove that reading.

NULL HALF-LIFE ACROSS 30 COMPARABLE FRONT-PAGE THREADS	MIN	P25	MEDIAN	P75	MAX
hours	5.76	6.88	7.69	8.80	16.94

withdrawn **The spread**

The incident’s thread is 7.05 h → **33rd percentile**, rank 10 of 30. It decays **more slowly** than the median thread. The number measures Hacker News, not the event. And a direct ratio bootstrap of lookups vs media gives 1.120, CI (0.788, 1.740) — contains 1.

survives **What remains**

Every front-page aggregator thread turns over in about eight hours — exactly as Wu & Huberman reported in 2007. So a conversion scorecard read at **48 hours** is always reading a channel that closed a day and a half earlier, whatever the event.

Every surviving number, re-tested against its strongest objection

OBJECTION	TEST	OUTCOME
Log-space OLS is biased	refit by raw-scale NLS	magnitudes move up to 2.7×; ordering invariant
Residuals autocorrelated	moving-block bootstrap	intervals widen 3–18%; all verdicts unchanged
Entity page contaminated	drop the Kimi-K3 release days	half-life 5.450 → 5.413 d (0.99×)
Breakpoint chosen by search	charge 2·ln C to AIC	two of three channels survive, not three
Null not season-matched	rebuild from prior-year Jul–Sep	68.2 → 53.8 ; matched figure reported
Comparison uses raw views	recompute in baseline-days	1.46× raw, 23.06× normalised; both reported
Overlapping intervals used as a test	bootstrap the ratio directly	(0.788, 1.740) — contains 1
Half-life may be platform-typical	null over 30 threads	33rd percentile — claim not supported

What this does not establish — and how it could be misused

Limitations

- **n = 1 event**, one matched control. The implications are hypotheses, not laws.
- **Proxy validity.** Pageviews measure lookups, not awareness. The entity page measures a **company**.
- **Small nulls.** 13 overlapping windows collapse to a handful of independent samples. 53.8 is a location, not a **p**-value.
- **Estimator dependence.** Half-lives move up to 2.7×; no single one is “the” quantity.
- **The rule is retrospective.** “Count the absorbing page” is an argmax over the outcome — a diagnostic, not an estimator.

Dual-use

The same finding that shows **attribution** drives attention could be read by a lab as an argument for **delaying or diffusing attribution** — and the counting result as an argument for steering coverage toward whichever page is least likely to register it.

We state it openly because it is **already** the incentive gradient labs face; because a public, checkable estimate is more useful to regulators, journalists and the breached party than to the one actor deciding whether to delay; and because it cuts symmetrically — it is equally an argument for a regime where attribution is timely **by requirement rather than by choice**.

No non-public data. No live system touched. No model capability involved. Nothing that lowers the cost of running an incident, and nothing that identifies a person.

What a communications team should do differently

1 · Declare the counting rule

Before asking how far an incident travelled, say whose page counts as “it”, in what units, against what null. For short windows some candidates — the event article — **don’t exist yet**.

2 · Don’t read the bubble’s clock

Every front-page thread closes in ~8 hours. A scorecard at 48 hours measures a channel that shut a day and a half earlier. **That’s a property of the medium.**

3 · Expect no vocabulary shift

The organisations moved; the concepts did not, to the limit of what we could detect. A **large** conversion is what the warning-shot argument needs, and there wasn’t one.

What survives is a **measurement discipline** rather than a story — state whose page you counted, in what units, against what null. Then most of what can be said about a warning shot’s reach becomes **checkable rather than arguable**.

\$0, no credentials, no network

```
# clone, install, check
pip install -e ".[dev]"

python -m pytest           # 124 tests, no network
python scripts/verify.py   # 121 checks vs committed results

# regenerate the analysis
python scripts/run.py all
python scripts/robustness_suite.py
python scripts/nulls_and_controls.py
python scripts/power_and_comparators.py
```

121/121

verify checks agree

124

tests, offline

3.5 MB

cached API responses shipped in-repo

MIT

code, data and figures

Fatimah Emad Eldin · Fatimah@trouve.works · Trouvé Works

Paper: paper/main.pdf — 8-page body, references and appendix excluded.