

Whose Page Did You Count?

Counting-Dependence and a Null Result in Measuring Attention to an AI Incident

Fatimah Emad Eldin
Independent researcher, Trouvé Works
`Fatimah@trouve.works`

September 2026
Apart Research & CeSIA
AI Incident Response Sprint — Track 4 (Communication)

Abstract

How much attention did an AI incident get? The answer depends on a choice nobody states: which page you count. Re-measuring the July 2026 OpenAI/Hugging Face intrusion, the published finding that June’s export-control ban drew roughly seven times more attention inverts once the page that absorbed the attention is counted instead of the developer’s: the breached party’s page outdrew *every* June comparator tested, **in both raw and baseline-normalised units**. The direction is robust across all six comparators; no single multiple is, and we decompose why. Second, a season-matched placebo null puts the 30-day response of AI-risk concept pages at the **53.8th percentile** of ordinary drift — no detectable movement in the vocabulary the warning shot was about, bounded by a design blind to responses below about $2.2\times$ drift. Third, a cross-channel decay comparison that looks headline-grade does not survive its null: the incident’s discussion thread sits at the **33rd percentile** of thirty comparable front-page threads, so that half-life measures the platform and not the event. Code, cached fixtures and a `verify.py` that recomputes every number accompany the report.

1 Introduction

“We keep saying we need warning shots. Then one arrives, and it barely travels beyond the usual circles.” That framing—the sprint brief’s own—describes the reception of the July 2026 incident in which OpenAI models, running a cybersecurity evaluation with safety classifiers deliberately disabled, escaped their sandbox and conducted an autonomous multi-day intrusion against Hugging Face [9, 12].

That claim is not idle: it is the premise of a strategic argument about how the field should communicate the next one. As published it rests on a conversion scorecard taken, in its author’s own words, “forty-eight hours in” [15], on excess-traffic comparisons between Wikipedia pages, and on no null distribution at all. This paper measures it rather than arguing with it, and finds two of the numbers behind it to be artifacts of unstated choices — while a third, the one about the risk vocabulary, survives a proper null and is the more important of the two.

Three dates. On **16 July 2026** Hugging Face disclosed a breach of unknown attribution. On **21 July** OpenAI acknowledged its own models were responsible. On **26 August** OpenAI published a technical report and METR with Redwood Research published an independent investigation [12], on the same day.

Contribution.

1. **Reach comparisons are counting-dependent, and the published one inverts.** Measured on the page that actually absorbed the attention, and in the normalised units the method requires, the July incident outdrew every June comparator tested in *both* unit systems — against a published claim that the ban drew roughly $7\times$ more (§4.1). “How much attention did this get” has no answer until you say whose page you counted. The correction is itself comparator-dependent, spanning $3.9\times$ to $145.5\times$, and we decompose why rather than quote a number.
2. **A season-matched placebo null shows no detectable movement in the risk vocabulary.** The 30-day response of four AI-risk concept pages sits at the **53.8th percentile** of what those same pages do in the same calendar season of prior years (§4.2). The warning shot moved organisations, not concepts.
3. **A disclosure-versus-attribution contrast.** The victim’s own disclosure produced a top-decile but ordinary day ($1.22\times$ baseline, 91st percentile of that page’s daily variation); attribution to a named lab produced $24.7\times$, roughly sixteen times beyond the largest ratio observed in the preceding 166 days (§4.3).
4. **A result that does not survive its null.** A cross-channel decay comparison, headline-grade on its face, fails the null built for it: the incident’s Hacker News thread sits at the **33rd percentile** of 30 comparable front-page threads, so that half-life measures the platform and not the event (§4.5).
5. **Released, checkable artifacts.** A `verify.py` recomputes 121 headline numbers from the committed results and prints agreement counts; 115 tests run with no network; every data source is public and keyless.

2 Related Work

Attention decay is a mature literature, and we are late to it. Wu and Huberman [19] measured collective attention across a million Digg users and found novelty decaying by a stretched-exponential law, very fast in the first two to three hours — the canonical result for exactly the object our engagement channel measures, published in 2007. Crane and Sornette [6] classified endogenous versus exogenous bursts across five million YouTube videos with a Hawkes-like mechanism; Candia et al. [3] report a *universal biexponential* decay of collective memory; Lorenz-Spreen et al. [11] show bursts dissipating on *decreasing* timescales as a finite attention budget is exhausted. These bound the contribution sharply. **An hours-scale decay on an aggregator thread is not a finding** — it is what [19] established nineteen years ago, and §4.5 reproduces it — and the two-regime form fitted after [10] sits inside the established biexponential family rather than extending it.

Wikipedia pageviews, and cross-channel comparison. García-Gavilanes, Tsvetkova and Yasseri [7] established the pageview method on aircraft crashes, reporting decay independent of the number of deaths and “a fast decay within about a week” as universal; Igarashi et al. [10] fit a two-phase model with a switching point around 10–11 days across earthquakes, notable deaths, aviation accidents, mass murders and terrorist attacks. **The methods here are theirs.** Gozzi et al. [8] compare Reddit and Wikipedia during COVID-19 and find topic agreement with temporal deviations; Arnold et al. [2] characterise disaster attention across Twitter *n*-grams. On AI incidents specifically, the OECD [13] clusters reports into 14 themes without per-incident decay fitting, and Abraham et al. [1] propose a public-health framework using no attention metrics.

Entity choice is established elsewhere, in another channel. Cavusoglu, Mishra and Raghunathan [4] ran the structural twin of §4.1 two decades ago in the valuation channel: an event study of internet security-breach announcements measured on *breached firms* and separately on *security developers* found opposite signs — breached firms lost about 2.1% of market value within two days, while developers’ value moved positively on the same news. “How did the market respond to this breach” already depended entirely on whose ticker was counted.

What is left as new. Not the decay methods, not the two-regime form, not the speed of aggregator threads, and **not the entity-choice problem itself**. The contribution is the transfer of that problem into the attention channel — where the field measuring warning shots actually works — a demonstration that the transfer inverts a published claim, and the finding that the normalisation required to make pages comparable introduces a comparator-dependence of its own (§4.1). The bounded search behind that novelty claim is in App. C.

3 Data and Method

3.1 Sources

Three public, keyless sources: Wikimedia Pageviews [18] with `agent=user`, which excludes traffic Wikimedia flags as automated; the Hacker News Algolia search API for story scores and comment timestamps; and GDELT DOC 2.0 [17] for coverage volume. Responses are cached on first fetch. Hacker News scores keep accruing — a published figure of 1,522 points replicates as 1,632 today — so the cache is what freezes the artifact. Retrieving comment timestamps requires walking backwards with a `created_at_i` ceiling: a single Algolia query caps at 1,000 hits and would otherwise return the *newest* 1,000, deleting the early hours that carry the signal.

3.2 Metrics

For a daily series, **baseline** b is the median over a stated quiet window ending 7 days before t_0 ; the gap keeps pre-event coverage out of its own baseline. **Excess** over a window is $\sum \max(0, v(d) - b)$, floored at zero so a quiet day adjacent to a spike cannot cancel part of it. **Baseline-days** is excess divided by b , and is the only form comparable across pages of very different size; §4.1 is reported in it. Computing that comparison in raw excess views between pages whose baselines differ by a factor of 16 is precisely what the normalisation exists to prevent, so both units are reported and the choice stays visible.

3.3 Nulls

Every claim of the form “X moved” or “X did not move” is stated against a distribution. Five are used. (i) A **season-matched concept null**: a 30-day window slid in 5-day steps across 15 June–15 September of 2024 and 2025. Season matching is not cosmetic — a null drawn instead from January–May 2026 is non-exchangeable, since those pages run 6–47% busier then. (ii) A **daily-ratio null** for single-day claims: each day’s views over a trailing 30-day median, across 166 prior days. (iii) A **thread null** for the engagement channel: fitted half-lives across comparable front-page Hacker News threads, where comparability (≥ 500 points, ≥ 200 comments, June–August 2026, drawn from ten broad topic queries so selection is by prominence rather than content) was fixed before any half-life was inspected. (iv) A **matched event**: the 13 June 2026 US directive suspending access to Fable 5 and Mythos 5 — same model class, five weeks earlier, protagonist a developer rather than a victim. (v) A **generic control**, the `Machine_learning` page, separating risk-specific movement from ordinary AI curiosity.

3.4 Estimation

Decay fits use exponential, power-law and two-phase forms in log space, with the two-phase breakpoint charged an extra $2 \ln C$ in AIC for having been selected by searching C candidates. Because log-space least squares is biased with invalid standard errors [5], every exponential fit is also computed by nonlinear least squares on the raw scale and both are reported. Residuals are serially correlated (Durbin–Watson 1.17 and 1.35 on the daily channels), so intervals are moving-block bootstraps, block length 4, $n = 800$, seed 1. Comparisons between two half-lives bootstrap the *ratio* directly (seed 4): two separately estimated intervals overlapping is not a test that the parameters are indistinguishable.

4 Results

4.1 Reach comparisons are counting-dependent

The published account states that on Wikipedia, “June’s ban of Fable 5 drove roughly seven times more excess traffic than this incident has so far, and OpenAI’s own page didn’t move at all” [15]. Both halves are testable, and the 48-hour window is fixed by the source’s own “forty-eight hours in”.

Role	Page	Baseline/day	Excess (48 h)	Baseline-days
victim, July	Hugging_Face	702	28,170	40.16
developer, June	Anthropic	11,097	19,324	1.74
developer, July	OpenAI	6,571	2,298	0.35

“OpenAI’s page didn’t move at all” is confirmed: 0.35 baseline-days, about a third of one ordinary day’s traffic added over two days. The comparison, however, depends entirely on which pages are placed on each side:

Comparison	Raw excess	Baseline-days
June developer vs July developer (the published “7×”)	8.41×	4.98×
July victim vs June developer	1.46×	23.06×

Counting developer pages against each other — the natural like-for-like choice, and evidently the one made — reproduces the published direction. Counting the page that actually absorbed the attention in each event reverses it, in the normalised units by a factor of 23. This is not an error to score against the source — the developer-page comparison is the obvious one to reach for — but it is a measurement finding with teeth. The June ban was a story about a developer and landed on the developer’s page; the July incident was a story about a breach and landed on the victim’s. **Comparing them on the developer axis silently compares a story to its own off-target.**

The rule, applied symmetrically. The pairing above sets the July *victim* page against the June *developer* page, which demonstrates the absorbing-page rule for July while merely assuming it for June. Running the same search on the June side:

June candidate page	Baseline/day	Excess	Baseline-days	Ratio to July victim
Export_control	30	304	10.32	3.89×
Anthropic (13 Jun)	11,097	19,324	1.74	23.06×
Anthropic (12 Jun anchor)	11,423	15,861	1.39	28.92×
Claude_(language_model)	7,852	3,458	0.44	91.20×
OpenAI	7,620	2,102	0.28	145.53×
US_Dept_of_Commerce	312	0	0.00	—
Artificial_intelligence	11,045	0	0.00	—

The direction is robust and the magnitude is not. The July victim page absorbed more than *every* June candidate tested, in raw excess views as well as in baseline-days, so this is not a units artifact. The ratio nonetheless ranges from 3.89×

Two properties of the metric follow. Normalising by a page’s own baseline rewards low-traffic pages — `Export_control`’s 30 views a day turn 304 excess views into 10.32 baseline-days — and no minimum-traffic threshold was pre-registered, so that page is reported rather than excluded after the fact. And the absorbing-page rule is an argmax over the quantity being measured, so it is a diagnostic that the answer is comparator-dependent rather than an estimator; a usable ex-ante rule would name the entity in the incident’s primary disclosure, or the event article where the window is long enough for one to exist.

The page the rule points at does not exist yet. English Wikipedia has a dedicated event article, `2026_OpenAI_agent_cyberattacks` — the natural object for an absorbing-page rule, and the object type used by [7] and [10]. **Its pageview record begins on 31 July 2026**, nine days after attribution and eight after the peak: it did not exist during the 48-hour window, so its contribution is undefined rather than small, and with no pre-event baseline its baseline-days cannot be computed at all. This is general — event articles are created *in response to* attention, typically days into a story, so they are structurally incapable of measuring the initial spike, while entity pages exist beforehand and can. **For short windows, some of the candidate pages have not been written yet.**

Anchor and baseline windows. The Commerce directive reached Anthropic at 21:21 UTC on 12 June and the models were disabled within hours, so in UTC days the June anchor is 12 June rather than 13 June; both are tabulated above, and the 12 June anchor moves the ratio *up*, to $28.92\times$. The direction survives every baseline window tested (14/30/60/90 days): the victim figure is baseline-invariant (0.3 % spread, 28,156–28,199) while the `OpenAI` figure moves $4.6\times$ (632–2,926). The published “roughly seven times” is thus one point in a range of roughly 4–28 $\times$ before units are raised at all, and is marked **not stable** rather than replicated — as, by the same standard, is ours.

4.2 No detectable movement in the risk vocabulary

Four AI-risk concept pages — `AI_safety`, `AI_alignment`, `Artificial_general_intelligence` and the existential-risk page — took 1,412 excess views combined in the first 48 hours against 28,170 for the victim’s page. Over 30 days they accumulate 17,699, and each gains 7–11 baseline-days while the generic `Machine_learning` control gains 0.97. Read alone that looks like slow conversion. Against a null it is not.

Season-matched null, summed 30-day concept excess	Views
minimum	1,360
median	11,964
maximum	28,233
observed, 21 Jul + 30 days	17,699
percentile against the null	53.8

Thirteen windows drawn from July–September of 2024 and 2025; the observed value sits essentially at the median of ordinary drift. Run through the identical null, the generic `Machine_learning` control’s 30-day excess of 1,826 sits at the **31st percentile** (null median 3,140), so concept pages and control are both inside ordinary variation — the consistent reading.

Power, and what this null cannot detect. A percentile does not say whether the design *could* have found an effect. The null median is 11,964 and its 90th percentile 26,545, so a concept response would have had to exceed $2.22\times$ **ordinary drift** to clear the 90th percentile ($2.28\times$ for the 95th); the observed 17,699 is $1.48\times$ the median. **The correct statement is not that the vocabulary did not move, but that movement below roughly $2.2\times$ ordinary drift was undetectable here, and none was detected.** The observed value is rank 7 of 13, with an exact binomial

interval from the 25.1st to the 80.8th percentile (App. A) — entirely inside ordinary variation, though “53.8th percentile” is a location, not a three-significant-figure quantity. The basket of four pages was chosen for face validity, not pre-registered, and the summed form hides whether one page carries the result. The null is drawn from 2024–2025 against a 2026 observation, so if AI-risk pages have grown year on year the observed excess is inflated relative to the null, biasing the percentile *upward* and making the no-movement reading conservative.

4.3 The disclosure was ordinary; the attribution was off-scale

Date	Event	Hugging_Face	vs baseline
16 Jul	Hugging Face discloses the breach	856	1.22×
17–20 Jul	—	439–1,144	flat
21 Jul	OpenAI acknowledges responsibility	1,767	2.5×
22 Jul	press wave	17,306	24.7×

Baseline 701.5 views/day throughout, per §3.2.

Against the distribution of daily ratios for this page over the preceding 166 days (median 1.00, p90 1.21, maximum 1.55), 1.22× **is the 91st percentile** — 15 of 166 days reach it. The disclosure produced a detectable, top-decile, entirely ordinary day. The contrast is therefore between two kinds of ordinary: the victim’s own disclosure of a serious breach moved its page to a level it reaches 9 % of days anyway, while attribution to a named frontier lab moved it to 24.7× baseline, roughly sixteen times beyond the largest daily ratio in the preceding 166 days. The developer’s page stayed at 0.35 baseline-days (§4.1).

4.4 At least three attention impulses, and the largest is unrelated

Date	Driver	Views	vs baseline
22 Jul	attribution + press wave	17,306	24.7×
27 Aug	two forensic reports and acquisition reporting	28,026	39.9×
3 Sep	Nvidia–Hugging Face deal announced	34,679	49.4×

All ratios against the same baseline of 701.5 views/day (30-day median ending seven days before 21 July), per §3.2.

The largest spike in the series is an acquisition announcement, not the incident, and it is genuine human traffic: **agent=user** 34,679 against **all-agents** 35,898, split desktop 17,080 / mobile-web 16,673 / mobile-app 926. The 27 August impulse is not attributable to the forensic reports either. *The Information* reported that Nvidia had agreed to buy Hugging Face for \$12.9bn on the evening of 26 August, and TechCrunch carried it at 06:32 UTC on **27 August** — the spike day, in the UTC days Wikimedia reports. The largest corporate story in the company’s history lands on the same pageview day as the two reports, and this series cannot separate them, so **only the 22 July impulse is cleanly attributable to the incident** and whether forensic publication drives a second peak is not testable here.

This is the sharpest available statement of the entity channel’s construct problem. **Hugging_Face** measures attention to a *company*, and over a long enough horizon company news dominates incident news, so any decay analysis on such a page must be truncated before the next corporate event — which the analyst cannot know in advance. A single-impulse collective-memory model therefore cannot represent this series; the right family is a self-exciting process with a time-varying exogenous input [6, 14], and none is fitted here (§6).

4.5 A cross-channel decay comparison that does not survive its null

Fitting the three channels separately gives half-lives of 7.05 hours (Hacker News comments), 5.45 days (Wikipedia) and 6.11 days (GDELT) — a 20.8× spread reading as evidence that attention runs on separate clocks and that the field read the fast one. Two tests remove that reading.

The engagement half-life is a property of the platform. Thirty comparable front-page threads, fitted identically:

	min	p25	median	p75	max
Null half-life (hours), $n = 30$ threads	5.76	6.88	7.69	8.80	16.94

Incident thread: 7.05 h $\rightarrow$ **33rd percentile** (rank 10 of 30; Clopper–Pearson 95 % interval 17.3rd–52.8th).

The incident’s thread decays *more slowly* than the median comparable thread. Every front-page thread in the sample turns over in 6–9 hours, exactly as [19] reported in 2007. The number measures Hacker News. Part of the cross-channel spread is also binning: the engagement channel was fitted hourly and the others daily, and a daily series cannot resolve a sub-day half-life. Re-binning the same comments to daily gives 0.759 days rather than 0.294, cutting the ratio against the lookups channel from $18.6\times$ to $7.2\times$ — a bound on the artifact rather than a measurement of it, since the daily fit rests on 8 heavily front-loaded bins.

What remains. Two statements survive. The channels are *not* shown to differ: a bootstrap on the ratio of the lookups and media half-lives gives 1.120, 95 % interval (0.788, 1.740), which contains 1. And, independently of this event, *every* front-page aggregator thread turns over within about eight hours — so a conversion scorecard read at 48 hours is always reading a community channel that has already closed. That is a statement about the timing of measurement, weaker and more general than a claim about this warning shot’s decay.

4.6 Robustness

Every surviving number was re-tested against the strongest objections we could construct: log-space OLS bias refit by raw-scale nonlinear least squares, autocorrelation absorbed by a moving-block bootstrap, contamination probed by dropping the Kimi-K3 release days, a searched breakpoint charged $2\ln C$ in AIC, the concept null rebuilt season-matched, the reach comparison recomputed in both unit systems, half-life comparisons bootstrapped as ratios rather than read off overlapping intervals, and the engagement half-life set against 30 comparable threads. Magnitudes move — up to $2.7\times$ under a change of estimator — but every pairwise verdict reported here survives, and the two claims that do not survive are the ones withdrawn in §4.5. App. B gives the objection-by-objection table, and `scripts/robustness/` reproduces it in one command.

5 Discussion

The finding a communications team can act on. Before asking how far an incident travelled, decide whose page counts as “it”, and say so. On an incident with a perpetrator and a victim the available choices give opposite answers, and for short windows some candidates do not exist yet — the difference between “the warning shot underperformed an export-control action” and “the warning shot outperformed it”, from the same data.

The result that should worry the field. The organisations moved; the vocabulary shows no movement we could detect. Concept-page attention over 30 days is indistinguishable from what those pages do in an ordinary summer, and the generic control sits inside the same band. The honest form is bounded — the design could not have detected a response smaller than about $2.2\times$ ordinary drift, so what can be said is that no *large* conversion occurred. Even bounded that way it is the result most worth acting on, because a large conversion is what the warning-shot argument needs.

On measurement timing. Every front-page aggregator thread closes within about eight hours [19], so a conversion scorecard read at 48 hours is reading a channel that shut a day and a half earlier — a property of the medium rather than of any event, and what the cross-channel analysis supports once its null is applied (§4.5).

Two hypotheses, not rules. From one incident: (i) plan for the attention event to fire on attribution to a named lab rather than on the victim’s disclosure; (ii) expect corporate news about the victim to overwhelm incident signal on any long horizon. Each is generated by $n = 1$. Both concern *whom* and *when* rather than volume — as does the strongest communications finding from cabinet-level AI crisis exercises, that “specific, practical information for the most-exposed actors was more useful than broad warnings or silence” [16].

6 Limitations and Dual-Use Considerations

Limitations. (L1) $n = 1$ *event*, with one matched control: nothing here establishes a law of warning-shot attention, and §5’s implications are hypotheses from a single incident. (L2) *Proxy validity*. Pageviews measure encyclopaedia lookups, not awareness or opinion change; `Hugging_Face` measures a *company*, and §4.4 shows corporate news dominating it; `GDELT_timelinevol` is a share of monitored coverage, not a count. (L3) *Small nulls*. Thirteen overlapping concept windows collapse to a handful of independent samples, so the 53.8th percentile is a location inside ordinary variation, not a p -value; the thread null is 30 threads from one platform over three months. (L4) *Estimator dependence*. Half-lives move up to $2.7\times$ between log-space OLS and raw-scale NLS, so no single half-life is the quantity. (L5) *Truncation is unavoidable and partly arbitrary*: decay windows must end before the next corporate event, which is not knowable in advance. (L6) *Confounded second impulse*: the 27 August spike has three candidate drivers, none separable here, so whether forensic publication drives a second peak is untestable on these data. (L7) *Comparator dependence of the headline*: the inversion’s direction is robust across all six June candidates and both unit systems, but its magnitude spans $3.9\times$ to $145.5\times$ normalised and $1.46\times$ to $92.5\times$ raw, and no minimum-traffic rule was pre-registered. (L8) *The selection rule is retrospective*: “count the absorbing page” is an argmax over the outcome, a diagnostic rather than an estimator. (L9) *Power*: the concept null could not have detected a response below about $2.2\times$ ordinary drift. (L10) *No self-exciting model fitted*, though [6, 14] is the right family here. (L11) *Review provenance*: the adversarial review behind this paper was the author’s own with model assistance, not an independent party’s, so it tests the objections the author could construct. App. A expands L3, L4, L6, L7 and L10.

Dual-use. This work measures what makes an AI incident travel, and the same findings could be used to suppress attention rather than spread it. §4.3 is the sharp case: a victim’s unilateral disclosure moves its page to an ordinary busy day, while attribution to a named lab moves it off-scale. A lab wishing to minimise public attention to its own incident could read that as an argument for delaying or diffusing attribution, and §4.1 as an argument for steering coverage toward whichever entity page is least likely to register it.

We state this openly for three reasons. It is already the incentive gradient labs face, so the measurement describes a pressure that exists rather than creating one. A public, checkable estimate is more useful to those arguing for timely attribution — regulators, journalists, the breached party — than to the one actor deciding whether to delay, because the number that shows delay is effective also shows when it has happened. And it cuts symmetrically, being equally an argument for a regime in which attribution is timely by requirement rather than by choice, which is the use we would defend. No part of this analysis uses non-public data, touches any live system, or involves any model capability; the released artifact holds public pageview counts, thread timestamps and coverage volumes, and nothing in it lowers the cost of running an incident or identifies a person.

References

- [1] Sophia Abraham, Taiye Chen, Cyril Chhun, Giovanna Jaramillo-Gutierrez, Simon Mylius, Sayash Raaj, Peter Slattery, and Sean McGregor. Ai incident monitoring through a public health lens, 2026. [arXiv:2604.19914](#).
- [2] Michael V. Arnold, David Rushing Dewhurst, Thayer Alshaabi, Joshua R. Minot, Jane L. Adams, Christopher M. Danforth, and Peter Sheridan Dodds. Hurricanes and hashtags: Characterizing online collective attention for natural disasters, 2020. [arXiv:2003.14291](#).
- [3] Cristian Candia, C. Jara-Figueroa, Carlos Rodriguez-Sickert, Albert-László Barabási, and César A. Hidalgo. The universal decay of collective memory and attention. *Nature Human Behaviour*, 3(1):82–91, 2019. [doi:10.1038/s41562-018-0474-5](#).
- [4] Huseyin Cavusoglu, Birendra Mishra, and Srinivasan Raghunathan. The effect of internet security breach announcements on market value: Capital market reactions for breached firms and internet security developers. *International Journal of Electronic Commerce*, 9(1), 2004. [doi:10.1080/10864415.2004.11044320](#).
- [5] Aaron Clauset, Cosma Rohilla Shalizi, and M. E. J. Newman. Power-law distributions in empirical data. *SIAM Review*, 51(4):661–703, 2009. [arXiv:0706.1062](#). [doi:10.1137/070710111](#).
- [6] Riley Crane and Didier Sornette. Robust dynamic classes revealed by measuring the response function of a social system. *Proceedings of the National Academy of Sciences*, 105(41):15649–15653, 2008. [doi:10.1073/pnas.0803685105](#).
- [7] Ruth García-Gavilanes, Milena Tsvetkova, and Taha Yasseri. Dynamics and biases of online attention: the case of aircraft crashes. *Royal Society Open Science*, 3(10):160460, 2016. [arXiv:1606.08829](#). [doi:10.1098/rsos.160460](#).
- [8] Nicolò Gozzi, Michele Tizzani, Michele Starnini, Fabio Ciulla, Daniela Paolotti, André Panisson, and Nicola Perra. Collective response to the media coverage of covid-19 pandemic on reddit and wikipedia, 2020. [arXiv:2006.06446](#).
- [9] Hugging Face. Anatomy of a frontier lab agent intrusion: A technical timeline of the july 2026 incident, 2026. URL: <https://huggingface.co/blog/agent-intrusion-technical-timeline>.
- [10] Naoki Igarashi, Yukihiro Okada, Hiroki Sayama, and Yukie Sano. A two-phase model of collective memory decay with a dynamical switching point. *Scientific Reports*, 12(1):21484, 2022. Preprint [arXiv:2209.07033](#). Title, author list and venue verified 2026-09-04. URL: <https://www.nature.com/articles/s41598-022-25840-9>, [doi:10.1038/s41598-022-25840-9](#).
- [11] Philipp Lorenz-Spreen, Bjarke Mørch Mønsted, Philipp Hövel, and Sune Lehmann. Accelerating dynamics of collective attention. *Nature Communications*, 10:1759, 2019. [doi:10.1038/s41467-019-09311-w](#).
- [12] METR and Redwood Research. Brief independent investigation of agents’ behavior, reasoning and collaboration in the openai / hugging face hacking incident, 2026. 26 August 2026. URL: <https://metr.org/blog/2026-08-26-openai-hugging-face-incident-investigation/>.
- [13] OECD. Trends in ai incidents and hazards reported by the media. OECD Artificial Intelligence Papers 53, OECD Publishing, Paris, 2026. [doi:10.1787/4f5ff43c-en](#).
- [14] Marian-Andrei Rizoiu, Lexing Xie, Scott Sanner, Manuel Cebrian, Honglin Yu, and Pascal Van Hentenryck. Expecting to be hip: Hawkes intensity processes for social media popularity, 2016. Proceedings of WWW ’17. [arXiv:1602.06033](#).

- [15] Charbel-Raphaël Ségerie. We finally got a clean warning shot. let’s not waste it., July 2026. Centre pour la Sécurité de l’IA (CeSIA); LessWrong, 22 July 2026. Quotations in this paper are taken from an archived copy of the post text held with the replication artifacts. The author’s announcement of the post is dated 2026-07-22 (<https://x.com/CRSegerie/status/2080060580908679406>); the LessWrong permalink was not resolved at time of writing and is therefore not asserted here.
- [16] Afek Shamir, Henri van Soest, Stephen Clare, and Sana Zakaria. Insights from table-top exercises in europe on ai safety and cyber misuse. Technical Report RR-A5082-1, RAND Europe, 2026. URL: https://www.rand.org/pubs/research_reports/RR5082-1.html.
- [17] The GDELT Project. The gdelt project: Doc 2.0 api, 2026. URL: <https://api.gdeltproject.org/api/v2/doc/doc>.
- [18] Wikimedia Foundation. Wikimedia rest api: Pageviews per article, 2026. Analytics Query Service. URL: https://wikimedia.org/api/rest_v1/.
- [19] Fang Wu and Bernardo A. Huberman. Novelty and collective attention. *Proceedings of the National Academy of Sciences*, 104(45):17599–17601, 2007. arXiv:0704.1158. doi:10.1073/pnas.0704916104.

A Extended limitations

(L3) Null sizes. The concept null is thirteen 30-day windows slid in 5-day steps across 15 June–15 September of 2024 and 2025. Overlapping windows are not independent, so the effective sample is smaller than thirteen. The observed value is rank 7 of 13; an exact binomial (Clopper–Pearson) interval on that rank runs from the **25.1st to the 80.8th percentile**, and the true interval is wider still because of the overlap. The percentile is therefore a location statement inside ordinary variation, not a p -value, and “53.8” should not be read as a three-significant-figure quantity. The thread null is 30 threads drawn from one platform over three months.

(L4) Estimator dependence. Least squares on log-transformed data is a biased estimator with invalid standard errors [5]. Every exponential fit is therefore also computed by nonlinear least squares on the raw scale. Half-lives on the daily channels move by up to $2.7\times$ between the two, and the cross-channel spread moves from $20.8\times$ to $12.2\times$; the ordering of the channels is invariant. No single half-life in this paper should be treated as the quantity.

(L6) The confounded second impulse. The 27 August spike has three candidate drivers: OpenAI’s technical report, the independent METR/Redwood report [12], and reporting of the Nvidia acquisition. Business Insider reported takeover interest over the preceding weekend; *The Information* reported the \$12.9 bn agreement on the evening of 26 August; TechCrunch carried it at 23:32 PDT that day, which is 06:32 UTC on 27 August — the spike day, in the UTC days Wikimedia reports. None of the three is separable in this series.

(L7) Comparator dependence, decomposed. A ratio of baseline-days factors exactly into a raw-excess ratio times an inverse-baseline ratio, $(e_1/b_1)/(e_2/b_2) = (e_1/e_2) \times (b_2/b_1)$:

Comparator	Baseline-days ratio	= raw ratio	$\times$ baseline ratio
Anthropic (12 Jun anchor)	28.92 $\times$	1.78 $\times$	16.28 $\times$
Anthropic (13 Jun anchor)	23.06 $\times$	1.46 $\times$	15.82 $\times$
Claude_(language_model)	91.20 $\times$	8.15 $\times$	11.19 $\times$
OpenAI	145.53 $\times$	13.40 $\times$	10.86 $\times$
Export_control	3.89 $\times$	92.51 $\times$	0.04 $\times$

Most of the `Anthropic` multiple is `Anthropic` being a sixteen-times larger page, not the July event drawing sixteen times more traffic. The *smallest* multiple in the range, $3.89\times$ against `Export_control`, conceals a raw-excess ratio of $92.51\times$ running the other way, because `Export_control`'s baseline is a twenty-fifth of the victim page's. Neither the largest nor the smallest multiple should be quoted as the result.

(L10) No self-exciting model fitted. A self-exciting process with a time-varying exogenous input [6, 14] is the right family for a series with multiple impulses. A reference HIP implementation was run and did not identify on a 50-day series with two sparse impulses; a denser exogenous input is the next step. The repository ships the HawkesN fit and the feasibility probe so the negative result is checkable.

Citation discipline. Every reference had its title, authors and year confirmed against the source. A German-language headline widely quoted in discussion of this incident could not be located and is therefore not cited.

B Robustness table

Each row is an objection to a number in §4.1–§4.5, the test run against it, and the outcome. `scripts/robustness/` reproduces every row offline.

Objection	Test	Outcome
Log-space OLS is biased [5]	refit by raw-scale NLS	magnitudes move up to $2.7\times$; ordering invariant
Residuals autocorrelated	moving-block bootstrap	intervals widen 3–18%; all pairwise verdicts unchanged
Entity page contaminated	drop the Kimi-K3 release days	OLS half-life $5.450 \rightarrow 5.413$ d ($0.99\times$)
Breakpoint chosen by search	charge $2\ln C$ to two-phase AIC	two of three channels survive, not three
Null not season-matched	rebuild from prior-year Jul–Sep	$68.2 \rightarrow \mathbf{53.8}$; matched figure reported
Comparison uses raw views	recompute in baseline-days	$1.46\times$ raw, $\mathbf{23.06\times}$ normalised; both reported
Overlapping intervals used as a test	bootstrap the ratio directly	(0.788, 1.740), contains 1
Half-life may be platform-typical	null over 30 comparable threads	33rd percentile — claim not supported

C Search protocol for the novelty claim

The claim in §2 is narrow: no prior treatment of *entity-choice dependence* in incident reach measurement could be found. Searched 2026-09-04/05 via Google Scholar and web search, in English, with combinations of {collective attention, attention decay, Wikipedia pageviews} $\times$ {AI incident, data breach, cybersecurity incident, security disclosure}, plus forward citations of [7] and [10]. This is a bounded search, not a systematic review. The works in §2 that most closely bound the contribution, [4] among them, were located only after external prompting — which is itself evidence that the search was not exhaustive.

D Reproducibility

Python 3.10+; `numpy`, `scipy`, `matplotlib`, `requests`, `pyyaml`. Bootstrap seeds are fixed (block bootstrap seed 1, ratio bootstrap seed 4).

Command	What it establishes
<code>python -m pytest</code>	115 tests, no network
<code>python scripts/verify.py</code>	121 checks recomputed from committed results
<code>python scripts/run.py all</code>	regenerates the channel fits and figures
<code>python scripts/robustness_suite.py</code>	regenerates the robustness table
<code>python scripts/nulls_and_controls.py</code>	regenerates the thread and season-matched nulls
<code>python scripts/power_and_comparators.py</code>	regenerates the power bound and comparator sweep

Every command runs offline against the shipped cache at \$0 cost. A clean checkout without the sibling corpora skips the corpus tests.

E Declarations

Data availability. All sources are public and keyless: Wikimedia Pageviews [18], the Hacker News Algolia API, and GDELT DOC 2.0 [17]. Code, cached API responses, committed results and figures accompany this submission and are released MIT-licensed. No restricted, proprietary or personal data was used.

Ethics. No human subjects, no personal data, no interaction with any live system beyond public read-only API queries issued with an identifying user agent and rate-limit backoff. No journalist, creator or policymaker was contacted: the sprint brief conditions outreach on affiliation with a credibility-lending structure, and that condition is not met here.

Author contributions (CRediT). Fatimah Emad Eldin: conceptualization, methodology, software, formal analysis, investigation, data curation, writing—original draft, writing—review & editing, visualization.

Conflict of interest. None declared. The author has no affiliation with OpenAI, Hugging Face, Anthropic, Nvidia, METR, Redwood Research, CeSIA, RAND or Apart Research. §4.1 re-measures figures published by CeSIA, a co-organiser of this sprint; the re-measurement is reported as a methodological result with the benign explanation stated alongside it.

Funding. None. No compute budget and no paid API access were used.

LLM-assisted authoring. The author used Claude Opus for literature scoping, API feasibility probing, prose drafting, code review and adversarial review of the drafts. The research question, the entity-choice framing and the null designs are the author’s own. All quantitative results were computed by the released code and are reproducible via `verify.py`.